

Table of Contents

1D Gaussians [10 pts]

Can the probability density function (pdf) of X ever take values greater than 1?

Write the expression for the pdf of a univariate gaussian:

Write the code for the function that computes the pdf at x .

Test your implementation against a standard implementation

What is the value of the pdf at $x = 0$? What is probability that $x = 0$?

A Write the transformation that takes $x \sim N(0, 1)$ to $z \sim N(\mu, \sigma^2)$

Write a code to sample from $N(\mu, \sigma^2)$

Test your implementation by computing statistics on the samples

Plot pdf and normalized histogram of samples

CSC 2506 Probabilistic Learning and Reasoning

Assignment I - Q1

Tianyi Liu liuti110 1005820827

1D Gaussians [10 pts]

Let X be a univariate random variable distributed according to a Gaussian distribution with mean μ and variance σ^2

Can the probability density function (pdf) of X ever take values greater than 1?

Answer: Yes. $\int f(x)dx = 1$ does not have any restrictions on the value of pdf for a specific $X = x$.

Write the expression for the pdf of a univariate gaussian:

Answer: $p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$

Write the code for the function that computes the pdf at x .

gaussian_pdf (generic function with 1 method)

```
• function gaussian_pdf(x; mean=0., variance=0.01)
•     #default variables mean and variance
•     #set with keyword arguments
•     return 1 / sqrt(2 * pi * variance) * e ^ (- (x - mean) ^ 2 / (2 * variance))
• end
```

Test your implementation against a standard implementation

E.g. from a library, e.g. Distributions.jl.

```
• using Test
```

```
• using Distributions: pdf, Normal
• # Note Normal uses N(mean, stddev) for parameters
```

Test.DefaultTestSet("Implementation of Gaussian pdf", Any[], 2, false)

```
• @testset "Implementation of Gaussian pdf" begin
•     x = randn()
•     @test gaussian_pdf(x) ≈ pdf.(Normal(0., sqrt(0.01)), x)
•     # ≈ is syntax sugar for isapprox, typed with '\approx <TAB>'
•     # or use the full function, like below
•     @test isapprox(gaussian_pdf(x, mean=10., variance=1), pdf.(Normal(10.,
• sqrt(1)), x))
• end
```

What is the value of the pdf at $x = 0$? What is probability that $x = 0$?

Answer: 3.989422804014327, in general $\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{\mu^2}{2\sigma^2}}$; $p(x = 0) = 0$.

A Write the transformation that takes $x \sim \mathcal{N}(0., 1.)$ to $z \sim \mathcal{N}(\mu, \sigma^2)$

A Gaussian with mean μ and variance σ^2 can be written as a simple transformation of the standard Gaussian with mean 0. and variance 1..

Answer: $\sigma x + \mu$

Write a code to sample from $\mathcal{N}(\mu, \sigma^2)$

Implement function returning n independent samples from $\mathcal{N}(\mu, \sigma^2)$ by transforming n samples from $\mathcal{N}(0., 1.)$

```
• function sample_gaussian(n; mean=0., variance=0.01)
•     # n samples from standard gaussian
•     x = randn(n)
•     # TODO: transform x to sample z from N(mean, variance)
•     z = sqrt(variance) .* x .+ mean
•     return z
• end;
```

Test your implementation by computing statistics on the samples

```
• using Statistics: mean, var
```

```
Test.DefaultTestSet("Numerically testing Gaussian Sample Statistics", Any[], 2, false)
```

```
• @testset "Numerically testing Gaussian Sample Statistics" begin
•     #TODO: choose some values of mean and variance to test
•     true_mean = 20.
•     true_var = 3.
•     #TODO: Sample 100000 samples with sample_gaussian
•     data = sample_gaussian(100000, mean=true_mean, variance=true_var)
•     #TODO: Use 'mean' and 'var' to compute statistics
•     stat_mean = mean(data)
•     stat_var = var(data)
•     #TODO: test statistics against true values
•     @test isapprox(stat_mean, true_mean, atol=1e-2)
•     @test isapprox(stat_var, true_var, atol=1e-2)
•     # hint: use isapprox with keyword argument atol=1e-2
• end
```

Plot pdf and normalized histogram of samples

Sample 10000 samples from a Gaussian with mean 10. and variance 2.0.

1. Plot the **normalized** histogram of these samples.
2. On the same axes plot! the pdf of this distribution.


Confirm that the histogram approximates the pdf.

(Note: with Plots.jl the function plot! will add to the existing axes.)

- using Plots

plot_hist_pdf (generic function with 4 methods)

- ```
function plot_hist_pdf(n=10000, mean=10., variance=2.)
 data = sample_gaussian(n, mean=mean, variance=variance)
 #histogram() #TODO
 histogram(data, norm=true, label="sample")
 #plot!() #TODO
 xs = mean - 3 * √variance : 0.01 : mean + 3 * √variance
 ys = gaussian_pdf.(xs, mean=mean, variance=variance)
 plot!(xs, ys, label="pdf", size=(800,600))
end
```



- plot\_hist\_pdf()

Answer: The histogram approximates the pdf.